

Coordination

Johan Montelius

KTH

2015

Why is coordination important?

Why is it a problem to implement?

1 / 29

2 / 29

Coordination

Failure detectors

Coordination in a distributed system:

- no fixed coordinator
- no shared memory
- failure of nodes and networks

The hardest problem is often knowing who is alive.

How do we detect that a process has crashed and how reliable can the result be?

- unreliable: result in *unsuspected* or *suspected* failure
- reliable: result in *unsuspected* or *failed*

Reliable detectors are only possible in synchronous systems.

3 / 29

4 / 29

- mutual exclusion - who is to enter a critical section
- leader election - who is to be the new leader
- group communication - same messages in the same order

- Safety: at most one process may be in critical section at a time
- Liveness: starvation free, deadlock free
- Ordering: enter in request happened-before order

5 / 29

6 / 29

Evaluation of algorithms

- Number of messages needed.
- Client delay: time to enter critical section
- Synchronization delay: time between exit and enter

A central server

Why not have one server that takes care of everything?

- *request a token* from the server
- *wait for a token* that grants access
- enter critical section
- exit critical section and *return the token*




Requirements: safety, liveness, ordering?

Evaluation: number of messages, client delay, synchronization delay

7 / 29

8 / 29

Pass a token around.



- pass a token around
- before entering the critical section - remove the token
- when leaving the critical section - release the token

Requirements: safety, liveness, ordering?

Evaluation: number of messages, client delay, synchronization delay

9 / 29

Why not complicate things?

To request entry:

- ask all other nodes for permission
- wait for all replies (save all requests from other nodes)
- enter the critical section
- leave the critical section (give permission to saved request)

otherwise:

- give permission to anyone

What could possibly go wrong?

10 / 29

Ricart and Agrawala

A request contains a *Lamport time stamp* and a *process identifier*.

Request can be ordered based on the time stamp and, if time stamps are equal, the process identifier.

When you're waiting for permissions and receive a request from another node:

- if the request is *smaller*, then give permission
- otherwise, save request

What order do we guarantee?

11 / 29

Maekawa's voting algorithm

Why ask all nodes for permission, why not settle for a *quorum*?

To request entry:

- ask all nodes your quorum for permission
- wait for all to vote for you:
 - queue requests from other nodes
- enter the critical section
- leave the critical section:
 - return all votes
 - vote for the first request if any in the queue


otherwise:

- if you have not voted:
 - vote for the first node to send a request
- if you have voted:
 - wait for your vote to return, queue requests from other nodes
 - when your vote is returned, vote for the first request if any in the queue

12 / 29

How do we form quorums?

- allow any majority of nodes
- divide nodes into groups, any two groups must share a node
- how small can the groups be?



All algorithms presented are more or less tolerant to failures.

Unreliable networks can be made reliable by retransmission (we must be careful to avoid duplication of messages)

Crashing nodes, even if we can detect them reliably, is a problem.

13 / 29

14 / 29

Election

Election, the problem of finding a leader in a group of nodes.

We assume that all nodes have unique identifiers.


Each node can *decide* which node to trust to be the *leader*.

Requirements:

- safety: if two nodes have decided they have decided to trust the same leader
- liveness: all nodes will eventually decide

Algorithms are evaluated on: number of messages and *turnaround time*.

A ring based approach



- a node starts an election
- the call is updated
- the leader is identified
- and proclaimed

Requirements: safety, liveness?

Evaluation: messages, turnaround?


15 / 29

16 / 29

Electing a new leader when the current leader has died.

- assumes we have *reliable failure detectors*
- all nodes know the nodes with higher priority

Assume we give priority to the nodes with lower process identifiers.



Requirements: safety, liveness?
turnaround?


Evaluation: messages,

17 / 29

18 / 29

Group communication

Multicast a message to *specified group of nodes*.



Reliability

- integrity: a message is only delivered once
- validity: a message is eventually delivered
- agreement: if a node delivers a message then all nodes will

Ordering of delivery:

- FIFO: in the order of the sender
- causal: in a happen-before order
- total: the same order for all nodes

Basic multicast

Assuming we have a reliable network layer this is simple.


A casted message is sent to all nodes in the group.

A received message is delivered.

What if nodes fail?

19 / 29


20 / 29



We have violated the *agreement requirement*.

How do we fix it?

We build a reliable multicaster



When receiving a message, forward it to all nodes.

Watch out for duplicates.

Alot of messages! $O(n^2)$

Reliable multicast often implemented by detecting failed nodes and then fix the problem.

In the assignment at the end we think it's working but it's not it's cheating!

We don't want to do that, that's too many messages —'.

21 / 29

22 / 29

Assume we first deliver a received message before we forward it.

Crashed nodes could have delivered a message.

Uniform agreement: if any node, correct or incorrect, delivers a message then all correct node will deliver the message.

Non-uniform agreement: if a correct node delivers a message then all correct node will


- FIFO: in the order of the sender
- causal: in a happen-before order
- total: the same order for all nodes

23 / 29

24 / 29

The simple way to implement ordered multicast.


- multicast the message to all nodes
- place in a hold-back queue
- multicast a *sequence number* to all nodes
- deliver in total order



Similar to Ricart and Agrawala

- multicast the message to all nodes
 - place in hold-back queue
 - propose a *sequence number*
 - select the highest
 - multicast the *sequence number* to all nodes
 - deliver in total order
- Why does this work?

unique identifier to break the tie



25 / 29

26 / 29

Causal ordering

Same as ISIS but, include a vector clock in the message.

Atomic Multicast

Atomic multicast: a reliable total order multicast.

Solves both leader election and mutual exclusion.

27 / 29

28 / 29

Coordination:

- mutual exclusion
- leader election
- group communication

Biggest problem is dealing with failing nodes.