

Problem

Distributed transactions

Johan Montelius

KTH

HT15


- Several independent transaction servers should be coordinated in one transaction.
- How do we coordinate operations to guarantee serial equivalence?

1 / 22

2 / 22

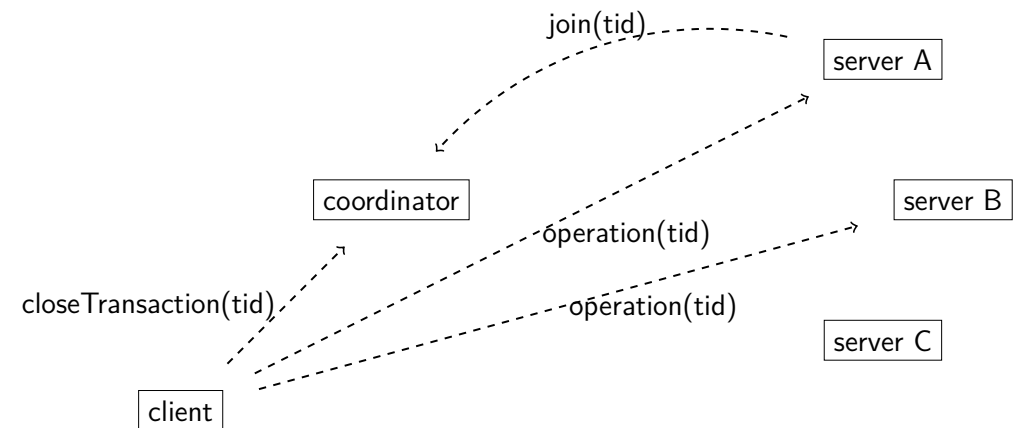
the architecture

transaction servers



3 / 22

transaction servers



4 / 22

one-phase commit

- Client sends closeTransaction to coordinator.
- Coordinator tells participants to commit the transaction.
- Problem: **If one of the server says I want to abort then that's the end because you need participation of the three servers actually the one-phase doesn't even exist!**
 - ?

two-phase commit

- phase one: ask participants to vote for commit or abort
 - if voting for commit one has to be able to commit even after a node crash
 - if anyone aborts all must abort
- phase two: inform all participants of the result

5 / 22

6 / 22

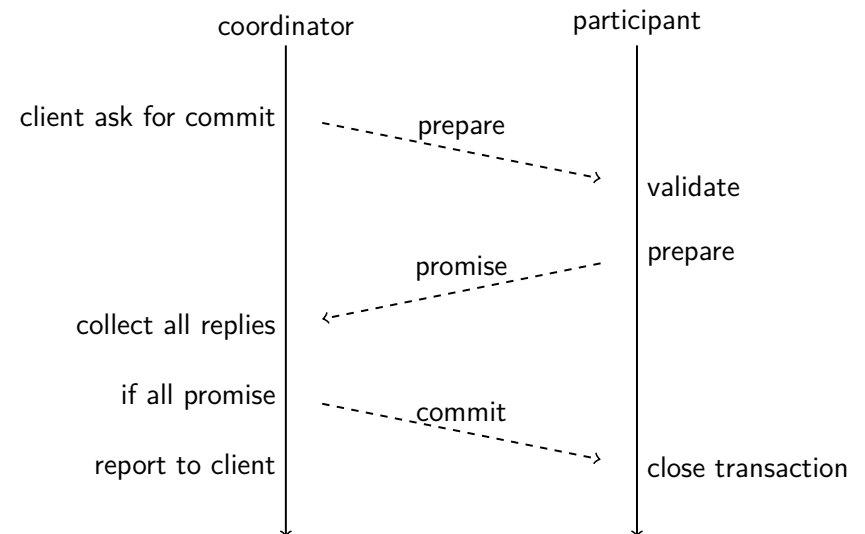
Consensus

Two-phase commit is a consensus protocol but:

- all servers must vote
- if any server wants to abort then we abort

7 / 22

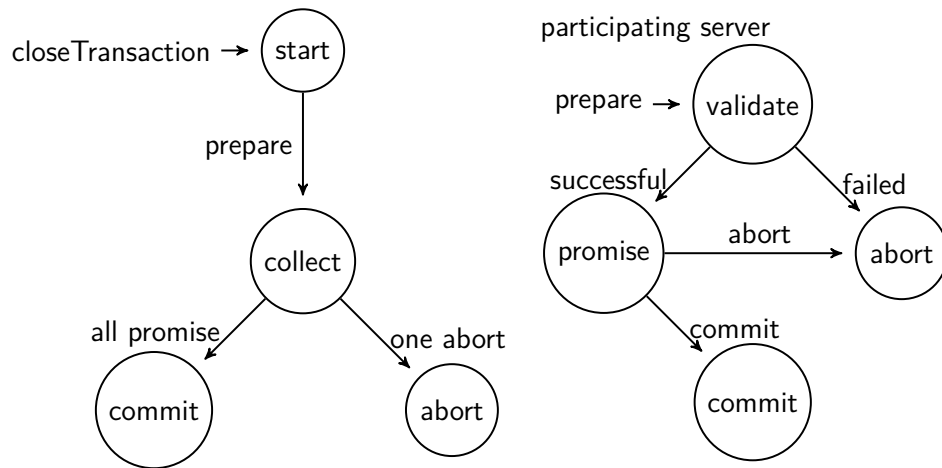
Two-phase commit



8 / 22

Two-phase commit

what if ..



It's very robust!

- a participating server crashes before making a promise
- a participating server crashes after having promised
- the coordinator crashes before asking for a promise
- the coordinator crashes but you have made a promise

two-phase commit can be suspended waiting for a crashed coordinator

9 / 22

10 / 22

if we know our peers


Three-phase commit

Assume that the participants know each other.

If the coordinator crashes:

- and no participant was told to commit, then it is safe to abort
- if one participant was told to commit, then we should all commit

What if the coordinator and one participant has crashed and none of the surviving participants have received a commit message?



- If in the *promised state* and coordinator crashes, and no non-crashed participant is in the *decided state* then abort, otherwise commit.
- If in the *decided state* and coordinator crashes then commit

Relies on perfect failure detectors - and that we know who is in the group.

11 / 22

12 / 22

It works but it is actually never used —


deadlock detection

What if:

- server A reports: S is waiting for T
- server B reports: T is waiting for U
- server C reports: U is waiting for S

Deadlock detected, let's do something

phantom deadlock



17 / 22

18 / 22

detection

How do we detect deadlocks?

With a marker ! you pass the marker and if you receive a marker it means that there is a circle then it means that there is a deadlock ! You abort and you break the circle ! penser à la petite mise en scène du prof !

optimistic concurrency control

Transactions should be validated in a total order.

What if transaction T is validated at A and transaction S at B?

19 / 22

20 / 22

A global timestamp that all transaction servers agree to.

Distributed transactions

- a global total order of transactions
- if one server needs to abort, then all should abort

Two-phase commit

- coordinator asks participants to prepare
- participants promise to commit (or aborts)
- coordinator directs participants to commit

Distributed deadlock

- hard to prevent
- simpler to detect

Concurrency control

- locks
- optimistic
- timestamp